

1. The `Caravan.csv` data set has information of 5822 individuals whom the Caravan insurance was offered. The variable `Purchase` (with categories Yes and No) indicates if an individual purchased the insurance. It can be verified that, in the dataset, only 348 individuals purchased the insurance policy. Thus, the data set is said to be *imbalanced*. And a strategy that predicts No always, would get a test error rate of 6% regardless of the predictor values. It seems that a predictive model should have a smaller error rate. But that is not important since the analyst should focus in a different rate, as explained below.

If the Caravan company tries to sell the insurance to any potential customer then the success rate would be 6%, which is not an effective marketing strategy. The company should rather focus on selling their insurance to customers who are likely to buy it. That is, to target their marketing effort in individuals who are predicted to buy it. Thus, the error rate of interest is the fraction of individuals who are correctly predicted to buy the insurance policy. We call this fraction the *precision*.

Consider for example that the confusion matrix that results from a classification model used with the Caravan data is

prediction	Actual Response	
	No	Yes
No	873	50
Yes	68	9

For this model the precision is $9/(68 + 9) = 0.117$ which is an improvement over 6%.

We want a model that results in the largest (test set) precision. Remove the variables `PVRAAUT` and `AVRAAUT` from the dataset. Create a test set consisting of the first 1,000 observations, and a train set with the remaining observations. Use `from sklearn.preprocessing import StandardScaler` if scaling is needed. Report the *test set precision* for the following models.

- a) A random forest model with 500 trees, `max_features = 29`, and `random_state = 0`.
 - b) A boosting model with `max_depth = 4`. Use 1000 trees, learning rate 0.01, and `random_state = 0`.
 - c) A multilayer perceptron with three hidden layers having 400 nodes each as follows

```
MLPClassifier(solver='lbfgs', random_state = 0,
              activation = 'tanh', hidden_layer_sizes=[400,400,400])
```
 - d) Logistic regression with the default value for `C`.
 - e) Repeat part (d) with a cut-off probability value equal to 0.25. `LogisticRegression` uses 0.50 as cut-off default value. It uses the cut-off value to predict the category. If the probability is larger than 0.5 the model predicts category Yes, otherwise it predicts No. Function `model.predict_proba()` may help to this end.
2. `scikit-learn` has a dataset of gray photos of public people. We want to build a classification model to predict the person's name from a new photo. Different individuals have different number of photos in the data. The following lines collect the photos of individuals with at least 60 photos in the dataset.

```
from sklearn.datasets import fetch_lfw_people
faces = fetch_lfw_people(min_faces_per_person = 60)
```

The size of each array is given by

```
print(faces.images.shape, faces.data.shape, faces.target.shape, faces.target_names.shape)
```

There are 1348 photos each 62×47 pixels large belonging to 8 different people (their names are in `faces.target_names.`). The 3D array `faces.images` stores the photos (1348 arrays of size 62×47). The array `faces.data` stores each photo in a row of length $62 \times 47 = 2914$. The features of the model are in `X = faces.data` and the response (target) variable in `y = faces.target`. Split the dataset into a train (75%) and a test set (25%). Use (`random_state = 42`). No need to scale the data.

- a) (10 pts.) Fit an SVC model as shown. Find the test accuracy rate.

```
svc = SVC(kernel='rbf', gamma = 0.01)
model = make_pipeline(svc)
model.fit(Xtrain, ytrain);
```

- b) (10 pts.) To improve the model accuracy we will transform the features using Principal Components, a topic we review in ISE 535. Use first 150 principal components as the new features of the model, as shown below, then report the test accuracy rate.

```
from sklearn.decomposition import PCA
pca = PCA(n_components=150, svd_solver='randomized',whiten=True,random_state=42)
svc = SVC(kernel='rbf', gamma = 0.01)
model = make_pipeline(pca,svc)
model.fit(Xtrain, ytrain);
```

- c) (10 pts.) Add argument `class_weight='balanced'` to the `SVC()` line above, run the code again and find the accuracy rate. Here we modify the loss function by weighting the loss of each row category by its class weight.

- d) (10 pts.) Use `GridSearchCV` to search best values for `C` and `gamma`. Try the grid values shown below. Report the test accuracy rate.

```
pca = PCA(n_components=150, svd_solver='randomized',whiten=True,random_state=42)
svc = SVC(kernel='rbf', class_weight='balanced')
model = make_pipeline(pca,svc)
param_grid = {'svc__C': [1, 5, 10],
              'svc__gamma': [0.0001, 0.0005, 0.001, 0.005]}
grid = GridSearchCV(model, param_grid,cv=5,iid = True)
grid.fit(Xtrain, ytrain)
model = grid.best_estimator_
```

- e) (10 pts.) Report the confusion matrix (for test set) of the best model in a dataframe with rows identifying the true name of the person and columns identifying the predicted name. What person has the largest error rate (that is, the person that is most difficult to predict)?

Your report must

- show your name and USC ID on the first page.
- be made of letter size pages in portrait format (not landscape).
- show Python commands fully displayed and not truncated.

Name your pdf file with your Name followed by your USC ID, for example `cesar_acosta - 0123456789.pdf`. Then submit your file onto Blackboard on the submission date.